

Aviraj Virape

Data Science Engineer (Intern)

Pune, Maharashtra | +91 99218 83944 | virapeaviraj@gmail.com

linkedin.com/in/avirajvirape | github.com/aviraj1805 | aviraj1805.github.io/portfolio | leetcode.com/u/aviraj_virape

SUMMARY

Final-year AI & Data Science student (CGPA 9.2), ex-founder (COEP I2I Startup Program), with research presented at an international conference. Builds end-to-end ML in Python: EDA, scikit-learn/XGBoost models, and Streamlit deployment; 275+ LeetCode problems.

EXPERIENCE

Data Science Intern | Innovexis

Feb 2026 – May 2026

- Built an ETL pipeline loading PhonePe Pulse JSON data (2018–2024) into a 9-table MySQL database.
- Developed a multi-page Streamlit and Plotly dashboard answering 10 business questions via SQL queries.

Founder | Mahaguru AI

Jan 2024 – Nov 2025

- Founded an AI student-guidance startup concept (career clarity, emotional support, learning plans); selected into COEP's I2I Startup Program for mentorship and incubation support.
- Designed a multi-agent LLM architecture (FastAPI, React) and curated a custom student–mentor dialogue dataset.
- Collected feedback from ~15 student users on early prototypes; presented the underlying research at INMEC-2026.

PROJECTS

Customer Churn Prediction – Kaggle Playground S6E3 | github.com/aviraj1805/Customer-Churn-Prediction | [Live](#)

- Trained XGBoost and LightGBM on ~594K telecom records (21 features) with 5-fold stratified cross-validation; XGBoost reached 0.916 mean ROC-AUC, with a preprocessing pipeline for encoding and median imputation.
- EDA identified the strongest churn drivers: electronic-check payment (49% churn), month-to-month contracts (42%), fiber-optic internet (41.5%), and first-year customers (49%). *Tech: Python, XGBoost, LightGBM, scikit-learn, Pandas, Seaborn*

CreditWise – Loan Approval Prediction | github.com/aviraj1805/CreditWise-Loan-Approval-System | [Live](#)

- Built a scikit-learn pipeline (imputation, scaling, one-hot encoding) on 950 applicant records with 18 features and 31% approvals, deployed as a Streamlit app with live predictions, confidence scores and feature importance.
- Trained a class-balanced Random Forest reaching 0.91 F1 and 0.98 ROC-AUC on a stratified 20% test set, versus 0.81 F1 for a logistic regression baseline; credit score and DTI ratio were the strongest predictors. *Tech: Python, scikit-learn, Pandas, NumPy, Streamlit*

IAMARS – Drone Detection and Tracking | <https://github.com/aviraj1805/intelligent-aerial-monitoring> | [Live](#)

- Led a team project fine-tuned YOLOv8n on the VisioDECT drone dataset with a leakage-aware block split, reaching 0.965 mAP@0.5 on a 1,800-image held-out test set.
- Pipeline runs at 88.5 FPS on an RTX 2050 GPU and 20 FPS on CPU; evaluated tracking with MOTA/IDF1 on 91 test videos, and intercept-point prediction in 3D simulation. Found and fixed a label-conversion bug that had inflated training mAP. *Tech: Python, Ultralytics YOLOv8, ByteTrack, NumPy, OpenCV*

SKILLS

Languages and Core: Python, SQL (MySQL), Data Structures and algorithms, Maths for AI

Machine Learning: scikit-learn, Supervised & Unsupervised Learning, Data Preprocessing, Data Analysis, EDA, Feature Engineering, Model Evaluation

Data & Databases: Pandas, NumPy, Matplotlib, Seaborn, Plotly

Tools & Deployment: Git, GitHub, Streamlit, Jupyter Notebook, Google Colab, Kaggle

PUBLICATION

“AI-Powered Reflective Mentorship” – Presented at the 1st International Conference on

Mar 2026

Interdisciplinary Advances in Mechanical, Electronics and Computer Engineering (INMEC-2026), Keystone School of Engineering, Pune

EDUCATION

B.E. Artificial Intelligence and Data Science | APCOER, Pune | CGPA 9.2

Expected May 2027

HSC (Science) | Yashodhara Junior College, Solapur | 78%

May 2022